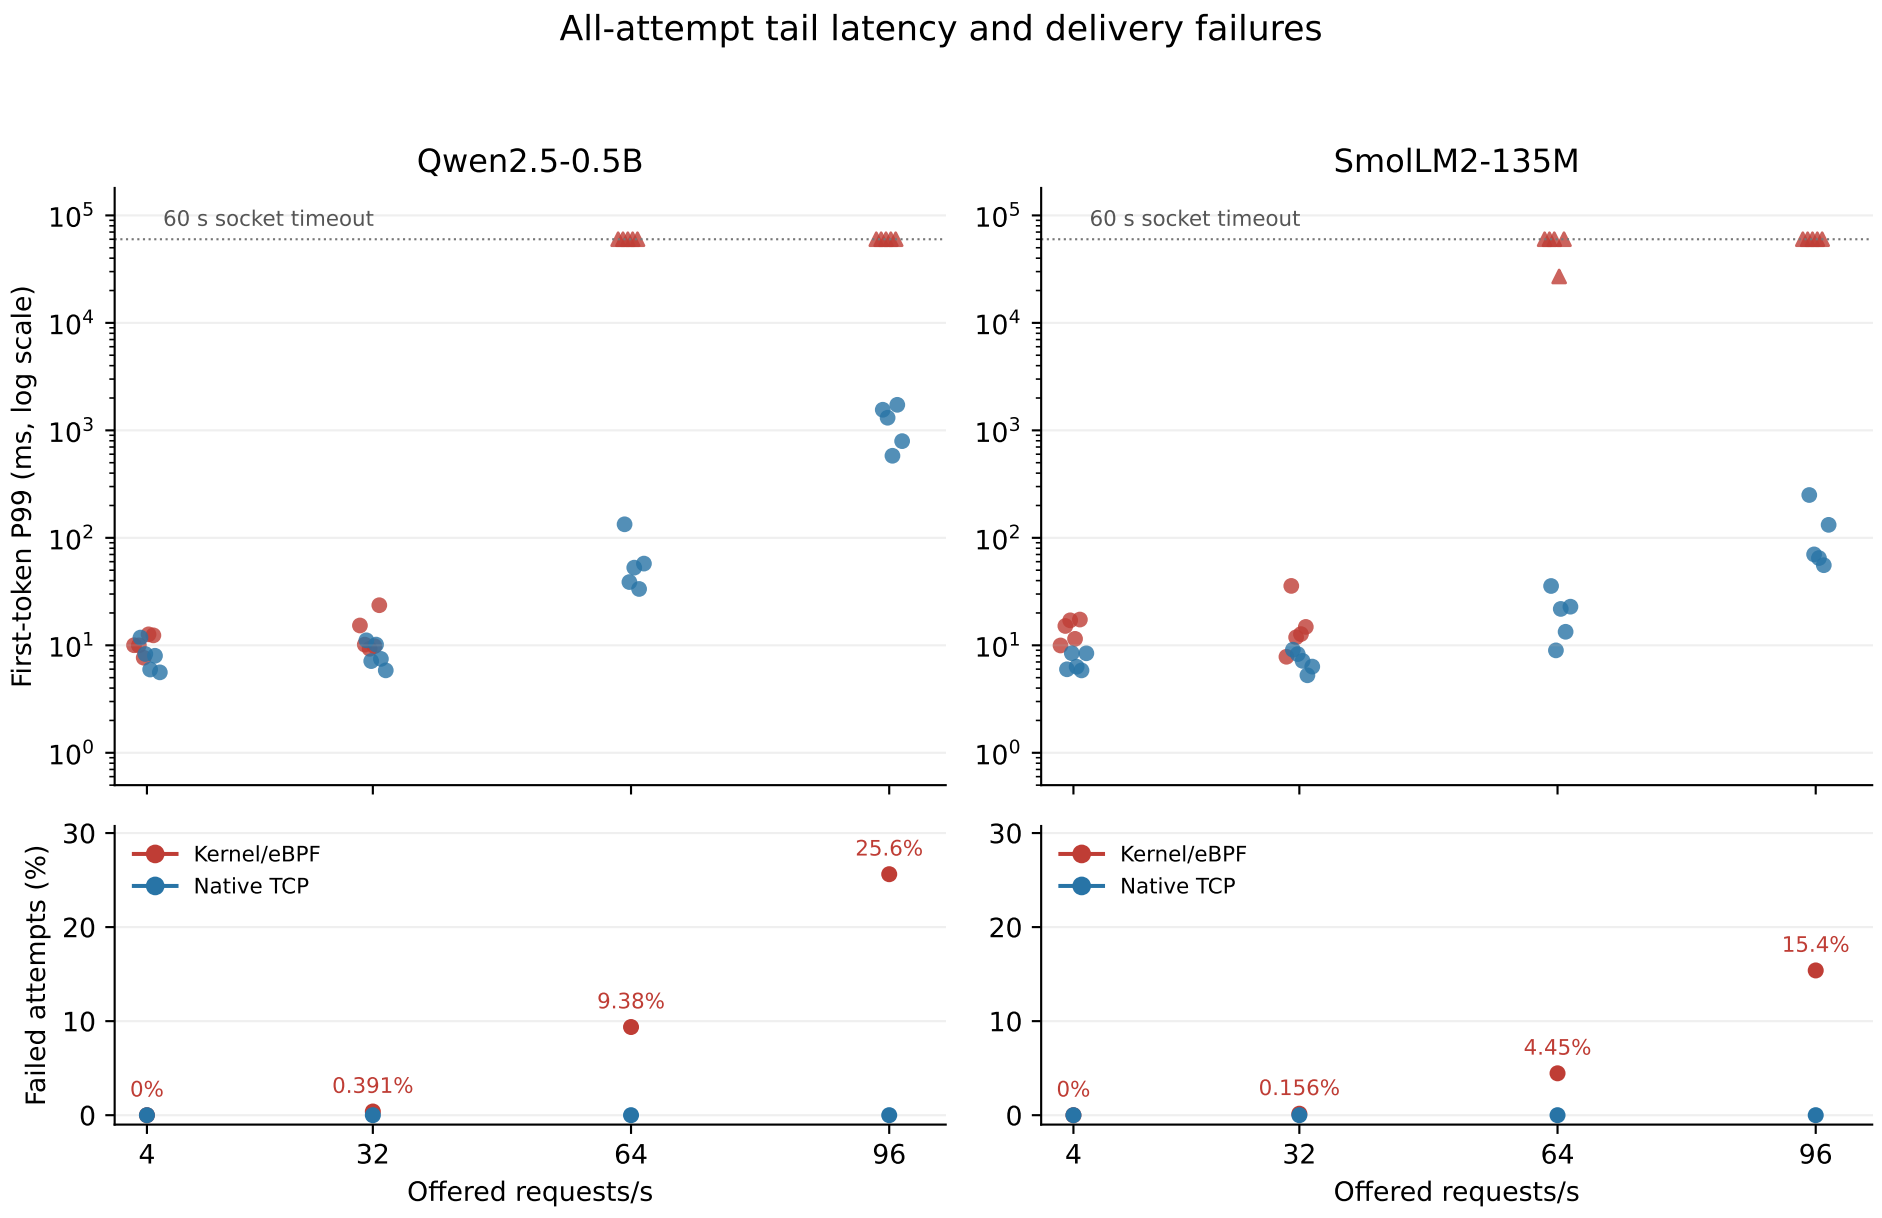
All-attempt tail latency and delivery failures
Qwen2.5-0.5B
SmolLM2-135M
60 s socket timeout
First-token P99 (ms, log scale)
Failed attempts (%)
Kernel/eBPF
Native TCP
25.6%
9.38%
0.391%
0%
15.4%
4.45%
0.156%
Offered requests/s
Five paired blocks/model; 256 attempts/cell; 64 output tokens. Each top point is one block.
Triangles show a lower bound with unknown upper tail, not a completed latency. No performance win is established.