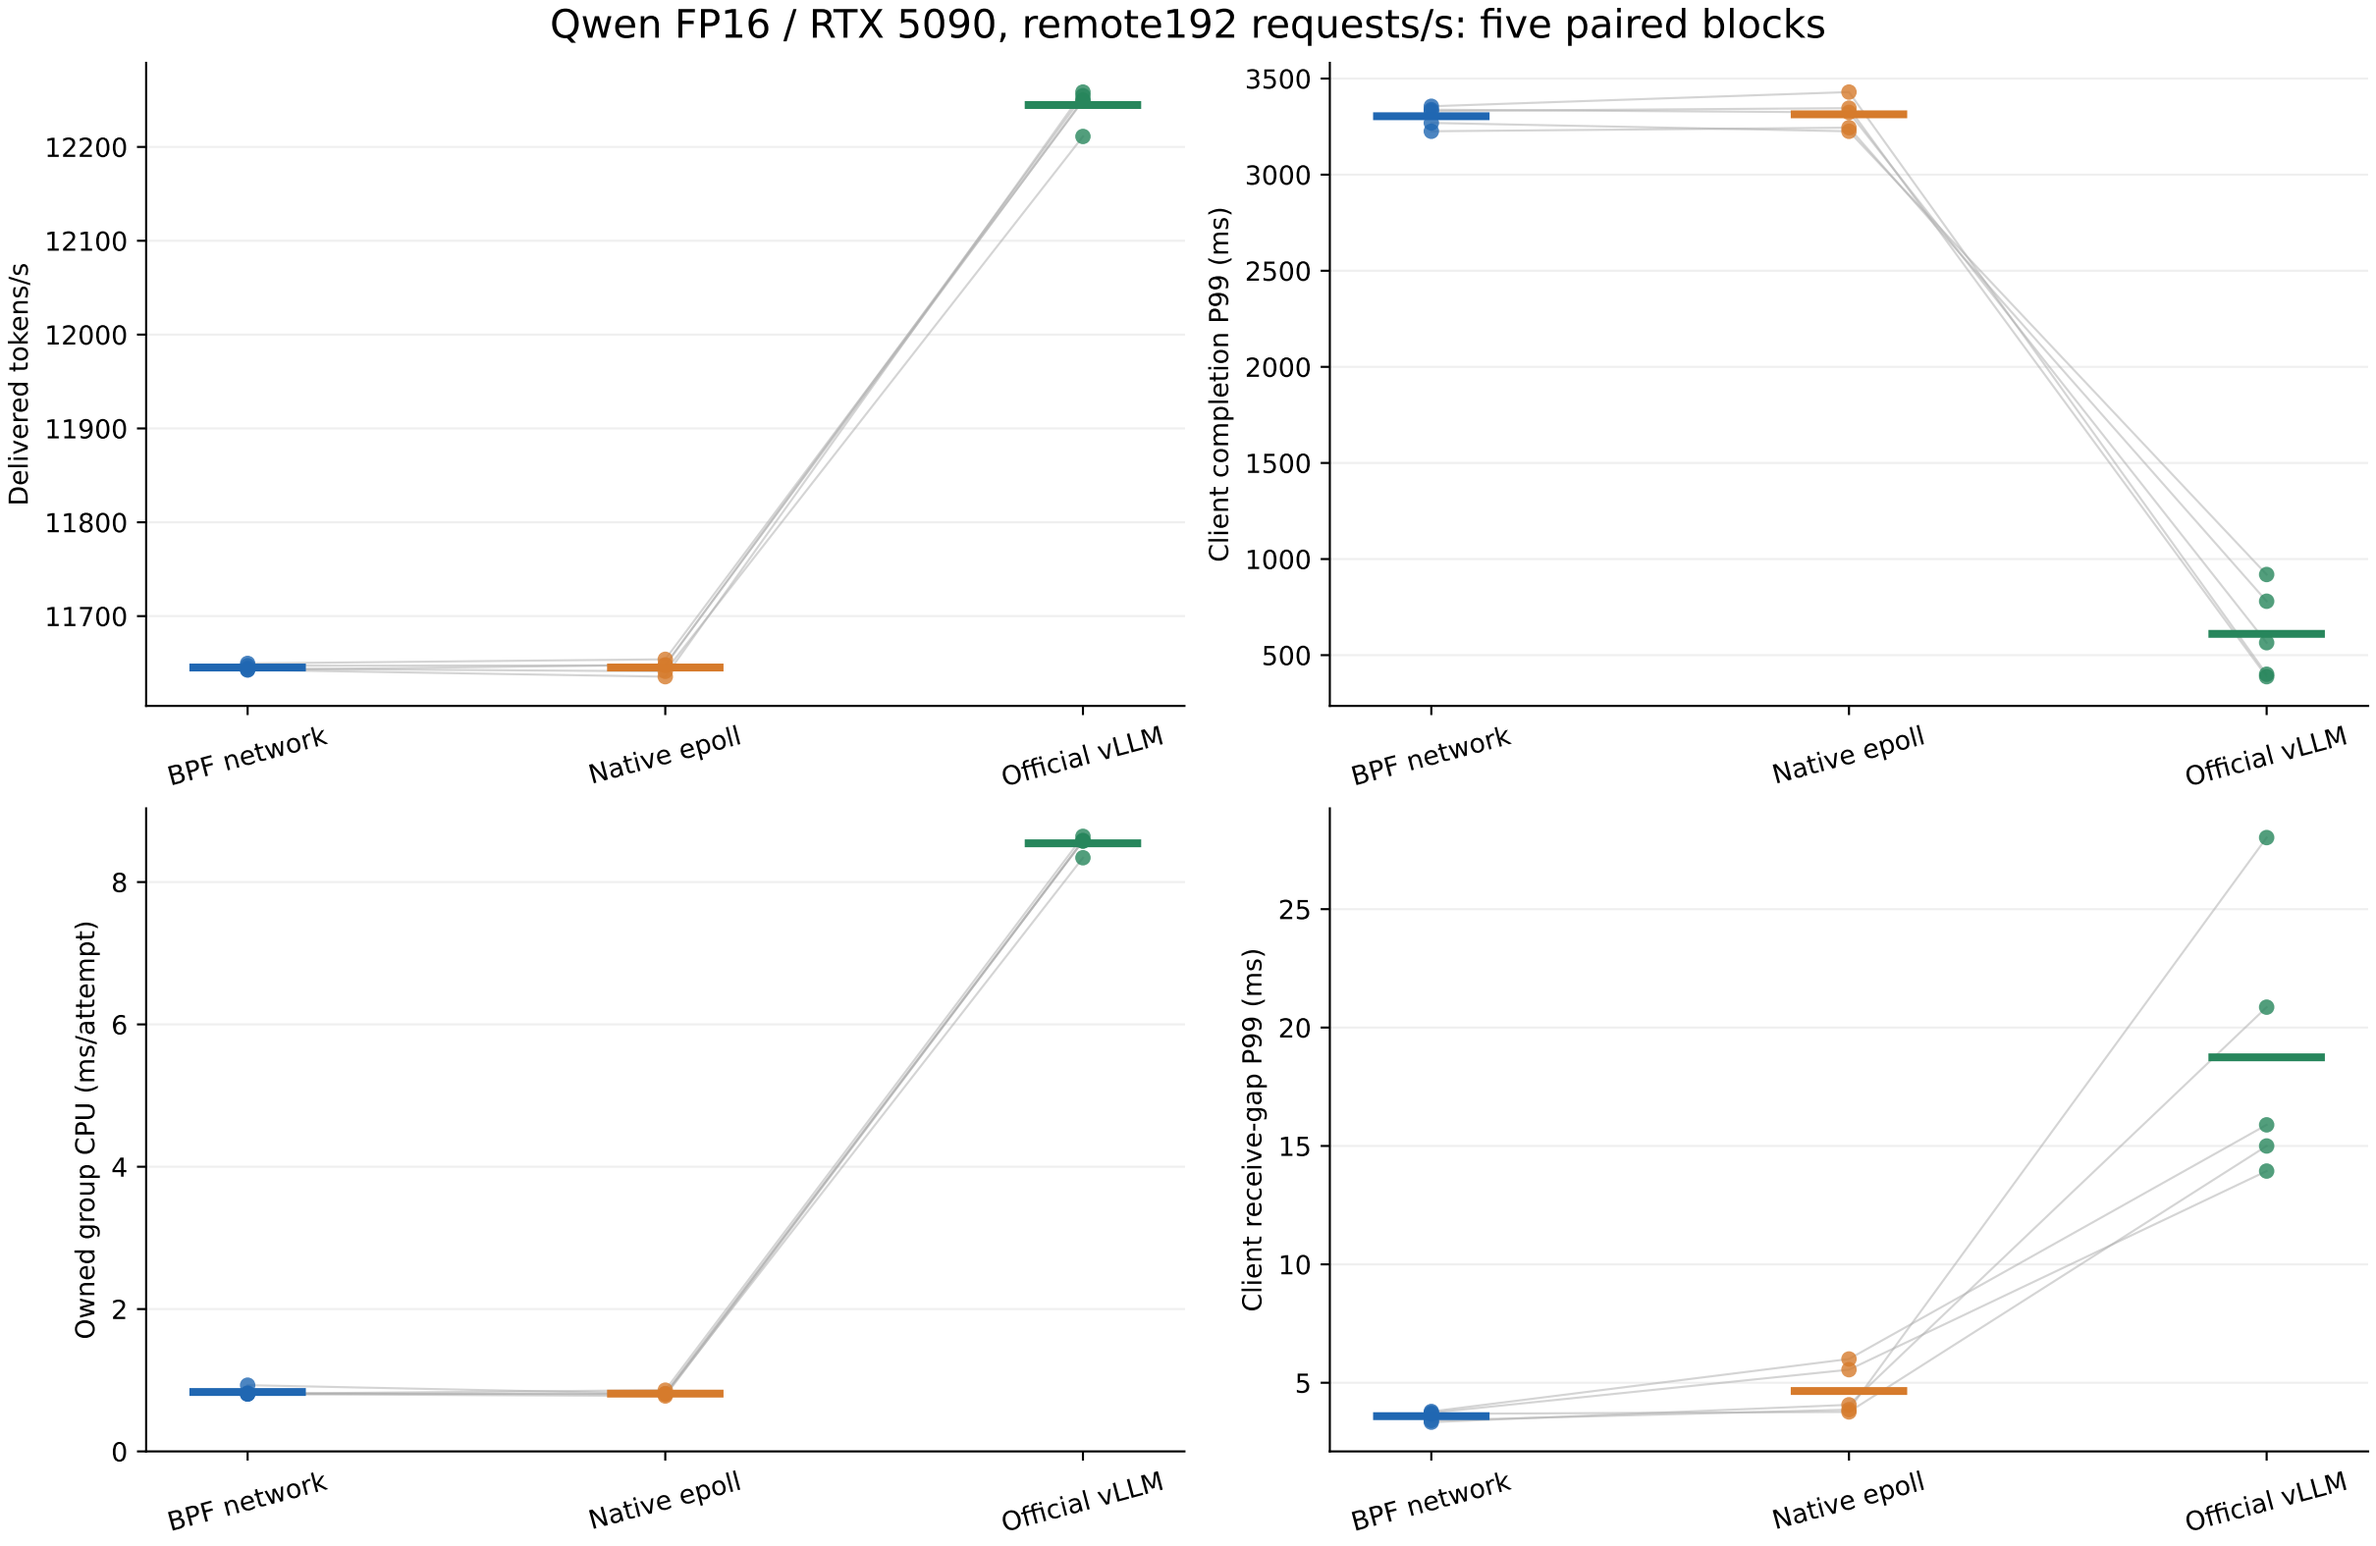
Qwen FP16 / RTX 5090, remote192 requests/s: five paired blocks
Delivered tokens/s
12200
12100
12000
11900
11800
11700
BPF network
Native epoll
Official vLLM
Client completion P99 (ms)
3500
3000
2500
2000
1500
1000
500
BPF network
Native epoll
Official vLLM
Owned group CPU (ms/attempt)
8
6
4
2
0
BPF network
Native epoll
Official vLLM
Client receive-gap P99 (ms)
25
20
15
10
5
BPF network
Native epoll
Official vLLM